

## **Reservoir** Labs

Power API  
2014.05.01

Sponsored by:

Defense Advanced Research Projects Agency

Microsystems Technology Office (MTO)

Program: Power Efficiency Revolution for Embedded Computing Technologies (PERFECT)

Issued by DARPA/CMO under Contract No: HR0011-12-C-0123

The views and conclusions contained in this document are those of the authors and should not be interpreted as representing the official policies, either expressly or implied, of the Defense Advanced Research Projects Agency or the U.S. Government.

Copyright © 2003-2014 Reservoir Labs, Inc.

# Contents

<b>1</b>	<b>Reservoir Labs Power API</b>	<b>1</b>
1.1	Overview . . . . .	1
1.2	High-Level Interface . . . . .	2
1.2.1	Hardware Behavior . . . . .	2
1.2.2	Speed Policy . . . . .	3
1.2.3	Scheduling Policy . . . . .	3
1.3	Low-Level Interface . . . . .	4
1.3.1	Speed Level . . . . .	4
1.3.2	Agility . . . . .	4
1.4	Example Code . . . . .	4
<b>2</b>	<b>Todo List</b>	<b>5</b>
<b>3</b>	<b>Module Documentation</b>	<b>7</b>
3.1	Limitations . . . . .	7
3.2	Speed Levels . . . . .	8
3.3	Speed Policies . . . . .	9
3.4	Scheduling Policies . . . . .	10
3.5	Unique Identifiers . . . . .	11
3.6	Units . . . . .	12
3.7	High-Level Interface Data Structures . . . . .	13
3.8	Error Codes . . . . .	14
3.9	High-Level Interface . . . . .	16

3.9.1	Function Documentation	16
3.9.1.1	pwr_change_hw_behavior	16
3.9.1.2	pwr_ecount_finalize	16
3.9.1.3	pwr_finalize	16
3.9.1.4	pwr_hw_behavior	17
3.9.1.5	pwr_initialize	17
3.9.1.6	pwr_is_initialized	17
3.10	Low-Level Interface	19
3.10.1	Function Documentation	19
3.10.1.1	pwr_agility	19
3.10.1.2	pwr_current_speed_level	20
3.10.1.3	pwr_energy_counter	20
3.10.1.4	pwr_islands	21
3.10.1.5	pwr_modify_speed_level	21
3.10.1.6	pwr_modify_voltage	22
3.10.1.7	pwr_num_islands	23
3.10.1.8	pwr_num_speed_levels	23
3.10.1.9	pwr_request_speed_level	24
<b>4</b>	<b>Data Structure Documentation</b>	<b>25</b>
4.1	e_data Struct Reference	25
4.1.1	Detailed Description	25
4.2	hw_behavior_t Struct Reference	25
4.2.1	Detailed Description	26
<b>5</b>	<b>File Documentation</b>	<b>27</b>
5.1	power_api.h File Reference	27
5.1.1	Detailed Description	31
5.2	power_api.h	31

## Chapter 1

# Reservoir Labs Power API

### 1.1 Overview

The Power API is divided into high-level and low-level interfaces. The high-level interface allows a programmer or compiler to specify power and energy management goals and choose strategies to achieve these goals. The implementation of strategies and the means to track and enforce power management goals is the responsibility of the implementor of the API on a platform.

The low-level interface of the Power API is close to the hardware and provides direct control over DVFS settings for voltage islands.

The high-level goals of the API are as follows:

- Provide a cross-platform interface for compilers and programmers to control power and energy consumption
- Be concise, intuitive and make minimal assumptions about the underlying hardware
- Maintain a level of abstraction high enough to allow implementation in terms of DARPA PERFECT team APIs.
- Take advantage of features provided by leading edge task-based runtime environments

The API assumes that any system components bound to the same voltage and frequency settings are grouped together in an island. An island is the atomic unit for which frequency and voltage can be modified through the Power API.

## 1.2 High-Level Interface

The high-level interface of the Power API is accessed through an initialization function and the data structures passed to this function.

The programmer (or compiler) must set up three things at Power API initialization:

- A model of hardware behavior
- A speed adjustment policy
- A scheduling policy

See also

[pwr\\_initialize\(\)](#)

Once configured, the combination of these 3 elements guides power and energy management decisions at program execution time. The 3 elements and associated data structures are described in the following sections. Power API implementations must define a default for each element on the targeted architecture.

### 1.2.1 Hardware Behavior

This defines the valid combinations of voltage and speed / speed level, possibly as functions of external factors such as the current temperature. Hardware behavior may be changed by hardware, software, or a combination of both.

See also

[hw\\_behavior\\_t](#)  
[pwr\\_hw\\_behavior\(\)](#)  
[pwr\\_change\\_hw\\_behavior\(\)](#)

**Todo** Clarify hardware behavior relationship with temperature / external factors.

Clarify definition of task, processing element.

Consider a near-threshold voltage architecture that may trade accuracy for power savings when voltage drops near threshold. An application that is resilient to errors would use a hardware behavior that allowed to use all voltage / frequency combinations supported by the architecture. An application that demands accurate results would use a hardware behavior that limited available voltage-frequency combinations to those well above threshold voltage.

### 1.2.2 Speed Policy

This defines a blanket policy for determining the speed level at which voltage islands are set.

Rather than exposing a notion of frequency, we expose a notion of speed level. This decision addresses the following:

- Permissible frequencies at discrete values
- Heterogeneity of architectures. Two different chips may have the same frequency but they won't have the same speed level (i.e., they won't execute the same piece of code in the same amount of time).

In the Power API, we define frequency, speed level, and speed:

- Frequency: The clock rate of a voltage island, in KHz
- Speed: A real number that corresponds to the absolute performance of a voltage island
- Speed Level: An integer greater than or equal to -1 that identifies a legal combination of voltage and frequency for a voltage island

#### See also

[Speed Policies](#)  
[pwr\\_request\\_speed\\_level\(\)](#)  
[pwr\\_modify\\_speed\\_level\(\)](#)  
[pwr\\_current\\_speed\\_level\(\)](#)  
[speed\\_policy\\_t](#)

### 1.2.3 Scheduling Policy

This defines a blanket policy for the run-time assignment of tasks to processing elements.

#### See also

[Scheduling Policies](#)  
[scheduling\\_policy\\_t](#)

**Todo** More options for scheduling policy. Specifically, space and time mapping.



## 1.3 Low-Level Interface

This interface allows the user to modify the speed level and voltage of an island. Speed level and voltage are not necessarily independent, so modifying one may modify the other.

A user is expected to be interested in upping the speed level or lowering the voltage knowing that the corresponding power consumption and speed will be negatively affected. An advanced user or compiler familiar with both island and application characteristics could modify speed and voltage in the same direction and still achieve power or performance gains. The goal of the low-level interface is to enable these modifications.

### 1.3.1 Speed Level

### 1.3.2 Agility

Agility is defined as the best and worst case amount of time it takes to switch from one speed level to another on a voltage island.

See also

[pwr\\_agility\(\)](#)

## 1.4 Example Code

```
#include <power_api.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char* argv[]) {
    retval_t retval;
    long num_islands;
    island_id_t islands;

    // Initialize with default values
    retval = pwr_initialize(NULL, NULL, NULL);

    // Get island ids
    retval = pwr_num_islands(&num_islands);
    islands = (island_id_t*)malloc(num_islands*sizeof(island_id_t));
    retval = pwr_islands(islands);

    // Set each island to max speed
    for (long i = 0; i < num_islands; ++i) {
        retval = pwr_request_speed_level(islands[i], PWR_MAX_SPEED_LEVEL);
    }

    // Finalize
    retval = pwr_finalize();
    return 0;
}
```

## Chapter 2

### Todo List

Global [pwr\\_agility](#) (const island\_id\_t island, const speed\_t from\_level, const speed\_t to\_level, agility\_t \*best\_case, agility\_t \*worst\_case)

Decide on unit for agility. Currently ns because of cpufreq default. If agility in cycles, at what speed level? Full speed? to\_level?

Global [pwr\\_modify\\_speed\\_level](#) (const island\_id\_t island, const int delta, const speed\_level\_t bottom)

Clarify behavior when an illegal speed level is requested

Clarify semantics of speed levels when multiple hardware behaviors are used

Global [pwr\\_modify\\_voltage](#) (const island\_id\_t island, const int delta)

Should we have a corresponding function to directly set voltage as we do with speed level?

Global [pwr\\_num\\_speed\\_levels](#) (const island\_id\_t island, long \*num\_speed\_levels)

Map speed level to integer between 0 and 100 inclusive?

page [Reservoir Labs Power API](#)

Clarify hardware behavior relationship with temperature / external factors.

Clarify definition of task, processing element.

More options for scheduling policy. Specifically, space and time mapping.



## Chapter 3

# Module Documentation

### 3.1 Limitations

#### Defines

- #define `PWR_MAX_ISLANDS` (1024L\*1024L\*1024L)  
*The maximum number of voltage islands in a system supported by the Power API.*
- #define `PWR_MAX_SPEED_LEVELS` (1024L\*1024L)  
*The maximum number of speed levels supported by the Power API.*

## 3.2 Speed Levels

### Defines

- #define `PWR_POWER_GATE` (-1)  
*Power to voltage island is off.*
- #define `PWR_CLOCK_GATE` ( 0)  
*Clock signal to voltage island is off.*
- #define `PWR_MIN_SPEED_LEVEL` (1)  
*Voltage island minimum speed level.*
- #define `PWR_MAX_SPEED_LEVEL` (LONG\_MAX)  
*Voltage island maximum speed level.*

## 3.3 Speed Policies

### Defines

- #define `PWR_SPEED_FIXED_MIN` (1)  
*Pin all voltage islands to minimum speed level.*
- #define `PWR_SPEED_FIXED_MAX` (2)  
*Pin all voltage islands to maximum speed level.*
- #define `PWR_SPEED_AS_NEEDED` (3)  
*Dynamically adjust speed level of voltage island to meet demand.*
- #define `PWR_SPEED_LOWLEVEL_API` (4)  
*Dynamically adjust speed level with Power API calls only.*

## 3.4 Scheduling Policies

### Defines

- #define `PWR_SCHED_DISTRIBUTE` (1)  
*Distribute tasks across processing elements such that hardware resource sharing between tasks is minimized.*
- #define `PWR_SCHED_CONCENTRATE` (2)  
*Distribute tasks across processing elements such that hardware resource sharing between tasks is maximized.*
- #define `PWR_SCHED_RANDOM` (3)  
*Distribute tasks across processing elements randomly.*
- #define `PWR_SCHED_EPS` (4)  
*Use energy proportional scheduling for space-time task mapping.*

## 3.5 Unique Identifiers

### Typedefs

- typedef long [island\\_id\\_t](#)  
*Integral unique identifier of a voltage island.*



## 3.6 Units

### Typedefs

- typedef double [voltage\\_t](#)  
*Voltage in volts.*
- typedef long [freq\\_t](#)  
*Frequency in KHz.*
- typedef long [speed\\_t](#)  
*Speed.*
- typedef long [speed\\_level\\_t](#)  
*Unit-less speed level.*
- typedef long [agility\\_t](#)  
*Agility in ns.*
- typedef long [power\\_t](#)  
*Power in watts.*
- typedef long [energy\\_t](#)  
*Energy in joules or microjoules (see function documentation)*
- typedef long [timestamp\\_t](#)  
*Time in seconds or nanoseconds (see function documentation)*

## 3.7 High-Level Interface Data Structures

### Data Structures

- struct `hw_behavior_t`  
*Defines legal (island\_id, voltage, frequency, speed, speed\_level) tuples.*

### Typedefs

- typedef long `speed_policy_t`  
*Defines policy for determining speed.*
- typedef long `scheduling_policy_t`  
*Defines policy for scheduling.*

## 3.8 Error Codes

### Defines

- #define `PWR_ARCH_UNSUPPORTED` (-3)  
*Feature unsupported by hardware.*
- #define `PWR_UNIMPLEMENTED` (-2)  
*Feature not implemented.*
- #define `PWR_UNINITIALIZED` (-1)  
*Power API has not been initialized.*
- #define `PWR_OK` (0)  
*Command executed successfully.*
- #define `PWR_ERR` (1)  
*Unspecified error.*
- #define `PWR_UNAVAILABLE` (2)  
*Feature temporarily unavailable.*
- #define `PWR_REQUEST_DENIED` (4)  
*Request was denied.*
- #define `PWR_INIT_ERR` (5)  
*Unspecified error during initialization.*
- #define `PWR_FINAL_ERR` (6)  
*Unspecified error during finalization.*
- #define `PWR_ALREADY_INITIALIZED` (7)  
*Attempt to initialize API after it has been initialized.*
- #define `PWR_IO_ERR` (8)  
*Unspecified input / output error.*
- #define `PWR_UNSUPPORTED_SPEED_LEVEL` (9)  
*Speed level not supported by hardware or API.*
- #define `PWR_UNSUPPORTED_VOLTAGE` (10)  
*Voltage not supported by hardware or API.*
- #define `PWR_ALREADY_MINMAX` (11)  
*Feature is already set to minimum or maximum value.*
- #define `PWR_OVER_E_BUDGET` (12)  
*Request denied, over energy budget.*
- #define `PWR_OVER_P_BUDGET` (13)  
*Request denied, over power budget.*
- #define `PWR_OVER_T_BUDGET` (14)  
*Request denied, over thermal budget.*
- #define `PWR_INVALID_ISLAND` (15)  
*Specified island does not exist.*

- #define `PWR_DVFS_ERR` (16)  
*Unspecified error when changing voltage and/or frequency.*
- #define `PWR_OVERFLOW_ERR` (17)  
*Requested value overflow.*

### Typedefs

- typedef int `retval_t`  
*Error codes returned by API functions.*

## 3.9 High-Level Interface

### Functions

- `retval_t pwr_is_initialized (int *is_initialized)`  
*Checks if the Power API has been initialized.*
- `retval_t pwr_initialize (const hw_behavior_t *hw_behavior, const speed_policy_t *speed_policy, const scheduling_policy_t *scheduling_policy)`  
*Allocates resources used by the Power API.*
- `retval_t pwr_finalize ()`  
*Frees resources used by the Power API.*
- `retval_t pwr_ecount_finalize ()`  
*This method is a quick fix to a strange problem caused by re-initializing and finalizing the ecount environment multiple times.*
- `retval_t pwr_hw_behavior (hw_behavior_t *hw_behavior)`  
*Retrieve the currently active hardware behavior.*
- `retval_t pwr_change_hw_behavior (hw_behavior_t *hw_behavior)`  
*Change currently active hardware behavior.*

### 3.9.1 Function Documentation

#### 3.9.1.1 `retval_t pwr_change_hw_behavior ( hw_behavior_t * hw_behavior )`

##### Parameters

<i>hw_behavior</i>	Pointer to new active hardware behavior struct
--------------------	--

##### Return values

<i>PWR_OK</i>	Active hardware behavior successfully changed
<i>PWR_UNINITIALIZED</i>	Power API has not been initialized thus hardware behavior cannot be changed

#### 3.9.1.2 `retval_t pwr_ecount_finalize ( )`

In particular, if the `ec_finalize` method is called in-between measurements then for some unknown reason measurements after `ec_finalize` always return zero.

#### 3.9.1.3 `retval_t pwr_finalize ( )`

Must be called after all Power API functions have been called.

## Return values

<i>PWR_OK</i>	All resources freed successfully
<i>PWR_UNINITIALIZED</i>	Power API has not been initialized thus cannot be finalized

3.9.1.4 `retval_t pwr_hw_behavior ( hw_behavior_t * hw_behavior )`

## Parameters

<i>hw_behavior</i>	Pointer to active hardware behavior struct
--------------------	--

## Return values

<i>PWR_OK</i>	Active hardware behavior returned successfully
<i>PWR_UNINITIALIZED</i>	Power API has not been initialized thus hardware behavior is unavailable

3.9.1.5 `retval_t pwr_initialize ( const hw_behavior_t * hw_behavior, const speed_policy_t * speed_policy, const scheduling_policy_t * scheduling_policy )`

Must be called before any other function in the Power API can be used.

## Parameters

<i>hw_behavior</i>	Relationship between voltage and speed on all voltage islands
<i>speed_policy</i>	Blanket policy for controlling voltage island speed levels
<i>scheduling_policy</i>	Blanket policy for runtime task scheduling

## Return values

<i>PWR_OK</i>	Initialized with no errors
<i>PWR_INIT_ERR</i>	Initialization failed
<i>PWR_ALREADY_INITIALIZED</i>	API is initialized

3.9.1.6 `retval_t pwr_is_initialized ( int * is_initialized )`

## Parameters

<i>is_initialized</i>	TRUE if initialized, FALSE otherwise
-----------------------	--------------------------------------

## Return values

<a href="#"><i>PWR_OK</i></a>	All cases
-------------------------------	-----------

## 3.10 Low-Level Interface

### Functions

- `retval_t pwr_num_islands` (long \*num\_islands)  
*The number of voltage islands controlled by the Power API.*
- `retval_t pwr_islands` (island\_id\_t \*islands)  
*Retrieve unique ID for all voltage islands.*
- `retval_t pwr_num_speed_levels` (const island\_id\_t island, long \*num\_speed\_levels)  
*Number of discrete speed levels, not including power and clock gated states, supported by a voltage island.*
- `retval_t pwr_current_speed_level` (const island\_id\_t island, speed\_level\_t \*current\_level)  
*The current speed level of a voltage island.*
- `retval_t pwr_request_speed_level` (const island\_id\_t island, const speed\_level\_t new\_level)  
*Requests speed level on the given voltage island.*
- `retval_t pwr_modify_speed_level` (const island\_id\_t island, const int delta, const speed\_level\_t bottom)  
*Requests a speed level modification on the given island.*
- `retval_t pwr_agility` (const island\_id\_t island, const speed\_t from\_level, const speed\_t to\_level, agility\_t \*best\_case, agility\_t \*worst\_case)  
*Calculates the cost of switching speed levels.*
- `retval_t pwr_modify_voltage` (const island\_id\_t island, const int delta)  
*Requests a voltage level modification of the given island.*
- `retval_t pwr_energy_counter` (const island\_id\_t island, energy\_t \*e\_j, energy\_t \*e\_uj, timestamp\_t \*t\_sec, timestamp\_t \*t\_nsec)  
*Value of monotonically increasing energy counter on the given island.*

### 3.10.1 Function Documentation

3.10.1.1 `retval_t pwr_agility ( const island_id_t island, const speed_t from_level, const speed_t to_level, agility_t * best_case, agility_t * worst_case )`

**Todo** Decide on unit for agility. Currently ns because of cpufreq default. If agility in cycles, at what speed level? Full speed? `to_level`?



## Parameters

<i>island</i>	The island of interest
<i>from_level</i>	Starting speed level
<i>to_level</i>	Finishing speed level
<i>best_case</i>	Minimum cost in ns
<i>worst_case</i>	Maximum cost in ns

## Return values

<a href="#"><i>PWR_OK</i></a>	Agility successfully returned
<a href="#"><i>PWR_ERR</i></a>	Agility not returned

3.10.1.2 `retval_t pwr_current_speed_level ( const island_id_t island, speed_level_t * current_level )`

## Parameters

<i>island</i>	The island to check speed level on
<i>current_level</i>	The current speed level of <i>island</i>

## Return values

<a href="#"><i>PWR_OK</i></a>	Speed level has been returned
<a href="#"><i>PWR_INVALID_ISLAND</i></a>	The given ID does not correspond to a voltage island.
<a href="#"><i>PWR_UNINITIALIZED</i></a>	Power API has not been initialized

3.10.1.3 `retval_t pwr_energy_counter ( const island_id_t island, energy_t * e_j, energy_t * e_uj, timestamp_t * t_sec, timestamp_t * t_nsec )`

Total energy consumed (J) since some undefined starting point is  $e_j + e_{uj} \cdot 10^{-6}$ . Elapsed time (sec) since some undefined starting point is  $t_{sec} + t_{nsec} \cdot 10^{-9}$ .

## Parameters

<i>island</i>	The island to measure energy consumption on
<i>e_j</i>	Energy consumed since some unspecified starting point on <i>island</i> , joule component
<i>e_uj</i>	Energy consumed since some undefined starting point on <i>island</i> , microjoule component
<i>t_sec</i>	Time since some unspecified starting point, seconds component
<i>t_nsec</i>	Time since some unspecified starting point, nanoseconds component

## Return values

<a href="#"><i>PWR_OK</i></a>	Energy counter successfully returned
<a href="#"><i>PWR_ERR</i></a>	Energy counter not returned
<a href="#"><i>PWR_OVERFLOW_ERR</i></a>	Energy counter overflowed, inaccurate results provided

3.10.1.4 `retval_t pwr_islands ( island_id_t * islands )`

Requires that `islands` points to at least `num_islands*sizeof(island_id_t)` bytes of memory.

## Parameters

<i>islands</i>	An array of <code>island_id_t</code>
----------------	--------------------------------------

## Return values

<a href="#"><i>PWR_OK</i></a>	Voltage island IDs have been returned
<a href="#"><i>PWR_UNINITIALIZED</i></a>	Power API has not been initialized

3.10.1.5 `retval_t pwr_modify_speed_level ( const island_id_t island, const int delta, const speed_level_t bottom )`

`delta` can be positive or negative.

## Parameters

<i>island</i>	The island to speed up or slow down
<i>delta</i>	The number of speed levels to increment by
<i>bottom</i>	The minimum legal speed level, can be one of <a href="#"><i>PWR_POWER_GATE</i></a> , <a href="#"><i>PWR_CLOCK_GATE</i></a> or <a href="#"><i>PWR_MIN_SPEED_LEVEL</i></a>

**Todo** Clarify behavior when an illegal speed level is requested

Clarify semantics of speed levels when multiple hardware behaviors are used

## Returns

[\*PWR\\_OK\*](#) Speed level was successfully modified

## Return values

<a href="#"><i>PWR_UNSUPPORTED_SPEED_LEVEL</i></a>	The requested speed level is not supported
<a href="#"><i>PWR_ALREADY_MINMAX</i></a>	Voltage island already at requested minimum or maximum speed level
<a href="#"><i>PWR_DVFS_ERR</i></a>	DVFS failed
<a href="#"><i>PWR_INVALID_ISLAND</i></a>	The given ID does not correspond to a voltage island.
<a href="#"><i>PWR_OVER_E_BUDGET</i></a>	Request denied, over energy budget (reserved for future use)
<a href="#"><i>PWR_OVER_P_BUDGET</i></a>	Request denied, over power budget (reserved for future use)
<a href="#"><i>PWR_OVER_T_BUDGET</i></a>	Request denied, over thermal budget (reserved for future use)
<a href="#"><i>PWR_UNINITIALIZED</i></a>	Power API has not been initialized

3.10.1.6 `retval_t pwr_modify_voltage ( const island_id_t island, const int delta )`

`delta` can be positive or negative.

**Todo** Should we have a corresponding function to directly set voltage as we do with speed level?

## Parameters

<i>island</i>	The island to change voltage on
<i>delta</i>	The number of voltage levels to modify by

## Return values

<a href="#"><i>PWR_OK</i></a>	Voltage level successfully modified
<a href="#"><i>PWR_ARCH_UNSUPPORTED</i></a>	Voltage level cannot be modified
<a href="#"><i>PWR_UNSUPPORTED_SPEED_LEVEL</i></a>	The requested speed level is not supported
<a href="#"><i>PWR_ALREADY_MINMAX</i></a>	Voltage island already at requested minimum or maximum voltage
<a href="#"><i>PWR_DVFS_ERR</i></a>	DVFS failed
<a href="#"><i>PWR_INVALID_ISLAND</i></a>	The given ID does not correspond to a voltage island.
<a href="#"><i>PWR_OVER_E_BUDGET</i></a>	Request denied, over energy budget (reserved for future use)

<i>PWR_OVER_P_BUDGET</i>	Request denied, over power budget (reserved for future use)
<i>PWR_OVER_T_BUDGET</i>	Request denied, over thermal budget (reserved for future use)
<i>PWR_UNINITIALIZED</i>	Power API has not been initialized
<i>PWR_ERR</i>	Voltage level not successfully modified

#### 3.10.1.7 `retval_t pwr_num_islands ( long * num_islands )`

##### Parameters

<i>num_islands</i>	Total voltage islands controlled by the Power API
--------------------	---

##### Return values

<i>PWR_OK</i>	Number of voltage islands has been returned
<i>PWR_UNINITIALIZED</i>	Power API has not been initialized

#### 3.10.1.8 `retval_t pwr_num_speed_levels ( const island_id_t island, long * num_speed_levels )`

The slowest speed level is '1' and speed levels increase monotonically until the fastest speed level at 'num\_speed\_levels'.

**Todo** Map speed level to integer between 0 and 100 inclusive?

##### Parameters

<i>island</i>	The ID of the island to get speed level count for
<i>num_speed_levels</i>	The number of speed levels supported by <i>island</i>

##### Return values

<i>PWR_OK</i>	Number of speed levels has been returned
<i>PWR_INVALID_ISLAND</i>	The given ID does not correspond to a voltage island.
<i>PWR_UNINITIALIZED</i>	Power API has not been initialized

### 3.10.1.9 `retval_t pwr_request_speed_level ( const island_id_t island, const speed_level_t new_level )`

Requires `-1 <= new_level <= num_speed_levels`. `'new_level == -1'` requests power gating. `'new_level == 0'` requests clock gating.

#### Parameters

<i>island</i>	The island to request speed level on
<i>new_level</i>	The requested speed level

#### Return values

<a href="#"><i>PWR_OK</i></a>	Speed level has been changed
<a href="#"><i>PWR_UNSUPPORTED_SPEED_LEVEL</i></a>	The requested speed level is not supported
<a href="#"><i>PWR_ALREADY_MINMAX</i></a>	Voltage island already at requested minimum or maximum speed level
<a href="#"><i>PWR_DVFS_ERR</i></a>	DVFS failed
<a href="#"><i>PWR_INVALID_ISLAND</i></a>	The given ID does not correspond to a voltage island.
<a href="#"><i>PWR_OVER_E_BUDGET</i></a>	Request denied, over energy budget (reserved for future use)
<a href="#"><i>PWR_OVER_P_BUDGET</i></a>	Request denied, over power budget (reserved for future use)
<a href="#"><i>PWR_OVER_T_BUDGET</i></a>	Request denied, over thermal budget (reserved for future use)
<a href="#"><i>PWR_UNINITIALIZED</i></a>	Power API has not been initialized

## Chapter 4

# Data Structure Documentation

### 4.1 e\_data Struct Reference

Definition of struct to hold energy consumption and timing information.

```
#include <ecount.h>
```

#### Data Fields

- double **start\_time\_ns**
- double **stop\_time\_ns**
- long long **values** [NUM\_EVENTS]
- char **names** [NUM\_EVENTS][EC\_MAX\_STR\_LEN]
- char **units** [NUM\_EVENTS][EC\_MIN\_STR\_LEN]

#### 4.1.1 Detailed Description

Definition at line [39](#) of file [ecount.h](#).

### 4.2 hw\_behavior\_t Struct Reference

Defines legal (island\_id, voltage, frequency, speed, speed\_level) tuples.

```
#include <power_api.h>
```

### Data Fields

- long [num\\_tuples](#)  
*Total number of (island\_id, voltage, frequency, speed, speed\_level) tuples.*
- [island\\_id\\_t](#) \* [island](#)  
*Island that this tuple applies to (required)*
- [voltage\\_t](#) \* [volts](#)  
*Voltage value of the tuple (required)*
- [freq\\_t](#) \* [freq](#)  
*Frequency value of the tuple (required)*
- [speed\\_t](#) \* [speed](#)  
*Speed value of the tuple (optional)*
- [speed\\_level\\_t](#) \* [speed\\_level](#)  
*Speed level value of the tuple (required)*

#### 4.2.1 Detailed Description

Definition at line [332](#) of file [power\\_api.h](#).

## Chapter 5

# File Documentation

### 5.1 power\_api.h File Reference

Copyright 2013-14 Reservoir Labs, Inc.

```
#include <limits.h>
```

#### Data Structures

- struct [hw\\_behavior\\_t](#)  
*Defines legal (island\_id, voltage, frequency, speed, speed\_level) tuples.*

#### Defines

- #define [PWR\\_MAX\\_ISLANDS](#) (1024L\*1024L\*1024L)  
*The maximum number of voltage islands in a system supported by the Power API.*
- #define [PWR\\_MAX\\_SPEED\\_LEVELS](#) (1024L\*1024L)  
*The maximum number of speed levels supported by the Power API.*
- #define [PWR\\_POWER\\_GATE](#) (-1)  
*Power to voltage island is off.*
- #define [PWR\\_CLOCK\\_GATE](#) ( 0)  
*Clock signal to voltage island is off.*
- #define [PWR\\_MIN\\_SPEED\\_LEVEL](#) (1)  
*Voltage island minimum speed level.*
- #define [PWR\\_MAX\\_SPEED\\_LEVEL](#) (LONG\_MAX)  
*Voltage island maximum speed level.*



- #define **PWR\_SPEED\_FIXED\_MIN** (1)  
*Pin all voltage islands to minimum speed level.*
- #define **PWR\_SPEED\_FIXED\_MAX** (2)  
*Pin all voltage islands to maximum speed level.*
- #define **PWR\_SPEED\_AS\_NEEDED** (3)  
*Dynamically adjust speed level of voltage island to meet demand.*
- #define **PWR\_SPEED\_LOWLEVEL\_API** (4)  
*Dynamically adjust speed level with Power API calls only.*
- #define **PWR\_SCHED\_DISTRIBUTE** (1)  
*Distribute tasks across processing elements such that hardware resource sharing between tasks is minimized.*
- #define **PWR\_SCHED\_CONCENTRATE** (2)  
*Distribute tasks across processing elements such that hardware resource sharing between tasks is maximized.*
- #define **PWR\_SCHED\_RANDOM** (3)  
*Distribute tasks across processing elements randomly.*
- #define **PWR\_SCHED\_EPS** (4)  
*Use energy proportional scheduling for space-time task mapping.*
- #define **PWR\_ARCH\_UNSUPPORTED** (-3)  
*Feature unsupported by hardware.*
- #define **PWR\_UNIMPLEMENTED** (-2)  
*Feature not implemented.*
- #define **PWR\_UNINITIALIZED** (-1)  
*Power API has not been initialized.*
- #define **PWR\_OK** (0)  
*Command executed successfully.*
- #define **PWR\_ERR** (1)  
*Unspecified error.*
- #define **PWR\_UNAVAILABLE** (2)  
*Feature temporarily unavailable.*
- #define **PWR\_REQUEST\_DENIED** (4)  
*Request was denied.*
- #define **PWR\_INIT\_ERR** (5)  
*Unspecified error during initialization.*
- #define **PWR\_FINAL\_ERR** (6)  
*Unspecified error during finalization.*
- #define **PWR\_ALREADY\_INITIALIZED** (7)  
*Attempt to initialize API after it has been initialized.*
- #define **PWR\_IO\_ERR** (8)  
*Unspecified input / output error.*

- #define `PWR_UNSUPPORTED_SPEED_LEVEL` (9)  
*Speed level not supported by hardware or API.*
- #define `PWR_UNSUPPORTED_VOLTAGE` (10)  
*Voltage not supported by hardware or API.*
- #define `PWR_ALREADY_MINMAX` (11)  
*Feature is already set to minimum or maximum value.*
- #define `PWR_OVER_E_BUDGET` (12)  
*Request denied, over energy budget.*
- #define `PWR_OVER_P_BUDGET` (13)  
*Request denied, over power budget.*
- #define `PWR_OVER_T_BUDGET` (14)  
*Request denied, over thermal budget.*
- #define `PWR_INVALID_ISLAND` (15)  
*Specified island does not exist.*
- #define `PWR_DVFS_ERR` (16)  
*Unspecified error when changing voltage and/or frequency.*
- #define `PWR_OVERFLOW_ERR` (17)  
*Requested value overflow.*

## Typedefs

- typedef long `island_id_t`  
*Integral unique identifier of a voltage island.*
- typedef double `voltage_t`  
*Voltage in volts.*
- typedef long `freq_t`  
*Frequency in KHz.*
- typedef long `speed_t`  
*Speed.*
- typedef long `speed_level_t`  
*Unit-less speed level.*
- typedef long `agility_t`  
*Agility in ns.*
- typedef long `power_t`  
*Power in watts.*
- typedef long `energy_t`  
*Energy in joules or microjoules (see function documentation)*
- typedef long `timestamp_t`  
*Time in seconds or nanoseconds (see function documentation)*

- typedef long [speed\\_policy\\_t](#)  
*Defines policy for determining speed.*
- typedef long [scheduling\\_policy\\_t](#)  
*Defines policy for scheduling.*
- typedef int [retval\\_t](#)  
*Error codes returned by API functions.*

## Functions

- [retval\\_t pwr\\_is\\_initialized](#) (int \*is\_initialized)  
*Checks if the Power API has been initialized.*
- [retval\\_t pwr\\_initialize](#) (const [hw\\_behavior\\_t](#) \*hw\_behavior, const [speed\\_policy\\_t](#) \*speed\_policy, const [scheduling\\_policy\\_t](#) \*scheduling\_policy)  
*Allocates resources used by the Power API.*
- [retval\\_t pwr\\_finalize](#) ()  
*Frees resources used by the Power API.*
- [retval\\_t pwr\\_ecount\\_finalize](#) ()  
*This method is a quick fix to a strange problem caused by re-initializing and finalizing the ecount environment multiple times.*
- [retval\\_t pwr\\_hw\\_behavior](#) ([hw\\_behavior\\_t](#) \*hw\_behavior)  
*Retrieve the currently active hardware behavior.*
- [retval\\_t pwr\\_change\\_hw\\_behavior](#) ([hw\\_behavior\\_t](#) \*hw\_behavior)  
*Change currently active hardware behavior.*
- [retval\\_t pwr\\_num\\_islands](#) (long \*num\_islands)  
*The number of voltage islands controlled by the Power API.*
- [retval\\_t pwr\\_islands](#) ([island\\_id\\_t](#) \*islands)  
*Retrieve unique ID for all voltage islands.*
- [retval\\_t pwr\\_num\\_speed\\_levels](#) (const [island\\_id\\_t](#) island, long \*num\_speed\_levels)  
*Number of discrete speed levels, not including power and clock gated states, supported by a voltage island.*
- [retval\\_t pwr\\_current\\_speed\\_level](#) (const [island\\_id\\_t](#) island, [speed\\_level\\_t](#) \*current\_level)  
*The current speed level of a voltage island.*
- [retval\\_t pwr\\_request\\_speed\\_level](#) (const [island\\_id\\_t](#) island, const [speed\\_level\\_t](#) new\_level)  
*Requests speed level on the given voltage island.*
- [retval\\_t pwr\\_modify\\_speed\\_level](#) (const [island\\_id\\_t](#) island, const int delta, const [speed\\_level\\_t](#) bottom)  
*Requests a speed level modification on the given island.*

- `retval_t pwr_agility` (const `island_id_t` island, const `speed_t` from\_level, const `speed_t` to\_level, `agility_t` \*best\_case, `agility_t` \*worst\_case)  
*Calculates the cost of switching speed levels.*
- `retval_t pwr_modify_voltage` (const `island_id_t` island, const int delta)  
*Requests a voltage level modification of the given island.*
- `retval_t pwr_energy_counter` (const `island_id_t` island, `energy_t` \*e\_j, `energy_t` \*e\_uj, `timestamp_t` \*t\_sec, `timestamp_t` \*t\_nsec)  
*Value of monotonically increasing energy counter on the given island.*

### 5.1.1 Detailed Description

Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License. Interface of the Power API for compiler-assisted energy-proportional scheduling.

Definition in file `power_api.h`.

## 5.2 power\_api.h

```

00001 /**
00002  * Copyright 2013-14 Reservoir Labs, Inc.
00003  *
00004  * Licensed under the Apache License, Version 2.0 (the "License");
00005  * you may not use this file except in compliance with the License.
00006  * You may obtain a copy of the License at
00007  *
00008  *     http://www.apache.org/licenses/LICENSE-2.0
00009  *
00010  * Unless required by applicable law or agreed to in writing, software
00011  * distributed under the License is distributed on an "AS IS" BASIS,
00012  * WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
00013  * See the License for the specific language governing permissions and
00014  * limitations under the License.
00015  */
00016
00017
00018
00019 /**
00020  * @file power_api.h Interface of the Power API for compiler-assisted
00021  * energy-proportional scheduling.
00022  *
00023  * \mainpage Reservoir Labs Power API
00024  *
00025  * \section over Overview
00026  * The Power API is divided into high-level and low-level interfaces. The
00027  * high-level interface allows a programmer or compiler to specify power and

```

```

00028 * energy management goals and choose strategies to achieve
00029 * these goals. The implementation of strategies and the means to track and
00030 * enforce power management goals is the responsibility of the implementor of
00031 * the API on a platform.
00032 *
00033 * The low-level interface of the Power API is close to the hardware and
00034 * provides direct control over DVFS settings for voltage islands.
00035 *
00036 * The high-level goals of the API are as follows:
00037 * - Provide a cross-platform interface for compilers and programmers to
00038 *   control power and energy consumption
00039 * - Be concise, intuitive and make minimal assumptions about the
00040 *   underlying hardware
00041 * - Maintain a level of abstraction high enough to allow implementation
00042 *   in terms of DARPA PERFECT team APIs.
00043 * - Take advantage of features provided by leading edge task-based runtime
00044 *   environments
00045 *
00046 * The API assumes that any system components bound to the same voltage and
00047 * frequency settings are grouped together in an island. An island is
00048 * the atomic unit for which frequency and voltage can be modified through
00049 * the Power API.
00050 *
00051 *
00052 * \section hl_desc High-Level Interface
00053 *
00054 * The high-level interface of the Power API is accessed through an
00055 * initialization function and the data structures passed to this function.
00056 *
00057 * The programmer (or compiler) must set up three things at Power API
00058 * initialization:
00059 *
00060 * - A model of hardware behavior
00061 * - A speed adjustment policy
00062 * - A scheduling policy
00063 *
00064 * @see pwr_initialize()
00065 *
00066 * Once configured, the combination of these 3 elements guides power and
00067 * energy management decisions at program execution time. The 3 elements
00068 * and associated data structures are described in the following sections.
00069 * Power API implementations must define a default for each element on the
00070 * targeted architecture.
00071 *
00072 *
00073 *
00074 * \subsection hw_behav Hardware Behavior
00075 *
00076 * This defines the valid combinations of voltage and speed / speed level,
00077 * possibly as functions of external factors such as the current temperature.
00078 * Hardware behavior may be changed by hardware, software, or a combination of
00079 * both.
00080 *
00081 * @see hw_behavior_t
00082 * @see pwr_hw_behavior()
00083 * @see pwr_change_hw_behavior()
00084 *
00085 * @todo Clarify hardware behavior relationship with temperature / external
00086 *   factors.
00087 *
00088 * Consider a near-threshold voltage architecture that may trade accuracy for
power
00089 * savings when voltage drops near threshold. An application that is
resilient
00090 * to errors would use a hardware behavior that allowed to use all voltage /
00091 * frequency combinations supported by the architecture. An application that
00092 * demands accurate results would use a hardware behavior that limited

```

```

available
00093 * voltage-frequency combinations to those well above threshold voltage.
00094 *
00095 *
00096 *
00097 * \subsection speed_policy Speed Policy
00098 *
00099 * This defines a blanket policy for determining the speed level at which
00100 * voltage islands are set.
00101 *
00102 * Rather than exposing a notion of frequency, we expose a notion of speed
00103 * level. This decision addresses the following:
00104 * - Permissible frequencies at discrete values
00105 * - Heterogeneity of architectures. Two different chips may have the same
00106 *   frequency but they won't have the same speed level (i.e., they won't
00107 *   execute the same piece of code in the same amount of time).
00108 *
00109 * In the Power API, we define frequency, speed level, and speed:
00110 *
00111 * - Frequency: The clock rate of a voltage island, in KHz
00112 * - Speed: A real number that corresponds to the absolute performance of a
00113 *   voltage island
00114 * - Speed Level: An integer greater than or equal to -1 that identifies a
00115 *   legal combination of voltage and frequency for a voltage
00116 *   island
00117 *
00118 * @see speed_policies
00119 * @see pwr_request_speed_level()
00120 * @see pwr_modify_speed_level()
00121 * @see pwr_current_speed_level()
00122 * @see speed_policy_t
00123 *
00124 *
00125 *
00126 * \subsection scheduling Scheduling Policy
00127 *
00128 * This defines a blanket policy for the run-time assignment of tasks to
00129 * processing elements.
00130 *
00131 * @see scheduling_policies
00132 * @see scheduling_policy_t
00133 *
00134 * @todo More options for scheduling policy. Specifically, space and time
00135 *   mapping.
00136 * @todo Clarify definition of task, processing element.
00137 *
00138 *
00139 *
00140 * \section ll_desc Low-Level Interface
00141 *
00142 * This interface allows the user to modify the speed level and voltage of an
00143 * island. Speed level and voltage are not necessarily independent, so
00144 * modifying one may modify the other.
00145 *
00146 * A user is expected to be interested in upping the speed level or lowering
00147 * the voltage knowing that the corresponding power consumption and speed will
00148 * be negatively affected. An advanced user or compiler familiar with both
00149 * island and application characteristics could modify speed and voltage in
00150 * the same direction and still achieve power or performance gains. The goal
00151 * of the low-level interface is to enable these modifications.
00152 *
00153 *
00154 * \subsection speedlevel Speed Level
00155 *
00156 * \subsection agility Agility
00157 * Agility is defined as the best and worst case amount of time it takes to
00158 * switch from one speed level to another on a voltage island.

```

```

00159  *
00160  * @see pwr_agility()
00161  *
00162  *
00163  *
00164  *
00165  *
00166  * \section example Example Code
00167  *
00168  * \code
00169  #include <power_api.h>
00170  #include <stdio.h>
00171  #include <stdlib.h>
00172
00173  int main(int argc, char* argv[]) {
00174      retval_t retval;
00175      long num_islands;
00176      island_id_t islands;
00177
00178      // Initialize with default values
00179      retval = pwr_initialize(NULL, NULL, NULL);
00180
00181      // Get island ids
00182      retval = pwr_num_islands(&num_islands);
00183      islands = (island_id_t*)malloc(num_islands*sizeof(island_id_t));
00184      retval = pwr_islands(islands);
00185
00186      // Set each island to max speed
00187      for (long i = 0; i < num_islands; ++i) {
00188          retval = pwr_request_speed_level(islands[i], PWR_MAX_SPEED_LEVEL);
00189      }
00190
00191      // Finalize
00192      retval = pwr_finalize();
00193      return 0;
00194  }
00195  \endcode
00196  *
00197  *
00198  *
00199  *
00200  *
00201  *
00202  *
00203  */
00204  #ifndef POWERAPI_H
00205  #define POWERAPI_H
00206  #ifdef __cplusplus
00207  extern "C" {
00208  #endif
00209
00210  #include <limits.h>
00211
00212  //=====
00213  // Constants
00214  //-----
00215  /** \addtogroup limits Limitations
00216   * @{
00217   */
00218
00219  /** The maximum number of voltage islands in a system supported by the Power
00220   API */
00221  #define PWR_MAX_ISLANDS (1024L*1024L*1024L)
00222  /** The maximum number of speed levels supported by the Power API */
00223  #define PWR_MAX_SPEED_LEVELS (1024L*1024L)
00224  /** @} */
00225

```

```

00226
00227
00228 /** \addtogroup speed_control Speed Levels
00229  * @{
00230  */
00231 /** Power to voltage island is off */
00232 #define PWR_POWER_GATE (-1)
00233
00234 /** Clock signal to voltage island is off */
00235 #define PWR_CLOCK_GATE ( 0)
00236
00237 /** Voltage island minimum speed level */
00238 #define PWR_MIN_SPEED_LEVEL (1)
00239
00240 /** Voltage island maximum speed level */
00241 #define PWR_MAX_SPEED_LEVEL (LONG_MAX)
00242 /** @} */
00243
00244
00245
00246 /** \addtogroup speed_policies Speed Policies
00247  * @{
00248  */
00249 /** Pin all voltage islands to minimum speed level */
00250 #define PWR_SPEED_FIXED_MIN (1)
00251
00252 /** Pin all voltage islands to maximum speed level */
00253 #define PWR_SPEED_FIXED_MAX (2)
00254
00255 /** Dynamically adjust speed level of voltage island to meet demand */
00256 #define PWR_SPEED_AS_NEEDED (3)
00257
00258 /** Dynamically adjust speed level with Power API calls only */
00259 #define PWR_SPEED_LOWLEVEL_API (4)
00260 /** @} */
00261
00262
00263
00264 /** \addtogroup scheduling_policies Scheduling Policies
00265  * @{
00266  */
00267 /** Distribute tasks across processing elements such that hardware resource
00268  sharing between tasks is minimized */
00269 #define PWR_SCHED_DISTRIBUTE (1)
00270
00271 /** Distribute tasks across processing elements such that hardware resource
00272  sharing between tasks is maximized */
00273 #define PWR_SCHED_CONCENTRATE (2)
00274
00275 /** Distribute tasks across processing elements randomly */
00276 #define PWR_SCHED_RANDOM (3)
00277
00278 /** Use energy proportional scheduling for space-time task mapping */
00279 #define PWR_SCHED_EPS (4)
00280 /** @} */
00281
00282 //=====
00283 // Typedefs
00284 //-----
00285
00286 /** \addtogroup id Unique Identifiers
00287  * @{
00288  */
00289 /** Integral unique identifier of a voltage island */
00290 typedef long island_id_t;
00291
00292 /** @} */
00293

```



```

00294 /** \addtogroup units Units
00295     * @{
00296     */
00297 /** Voltage in volts */
00298 typedef double voltage_t;
00299
00300 /** Frequency in KHz */
00301 typedef long freq_t;
00302
00303 /** Speed */
00304 typedef long speed_t;
00305
00306 /** Unit-less speed level */
00307 typedef long speed_level_t;
00308
00309 /** Agility in ns */
00310 typedef long agility_t;
00311
00312 /** Power in watts */
00313 typedef long power_t;
00314
00315 /** Energy in joules or microjoules (see function documentation) */
00316 typedef long energy_t;
00317
00318 /** Time in seconds or nanoseconds (see function documentation) */
00319 typedef long timestamp_t;
00320 /** @} */
00321
00322 //=====
00323 // High-level interface data structures
00324 //-----
00325 /** \addtogroup hli High-Level Interface Data Structures
00326     * @{
00327     */
00328
00329 /**
00330     * Defines legal (island_id, voltage, frequency, speed, speed_level) tuples
00331     */
00332 typedef struct hw_behavior {
00333     /** Total number of (island_id, voltage, frequency, speed, speed_level)
00334         tuples */
00335     long num_tuples;
00336     /** Island that this tuple applies to (required) */
00337     island_id_t island;
00338     /** Voltage value of the tuple (required) */
00339     voltage_t volts;
00340     /** Frequency value of the tuple (required) */
00341     freq_t freq;
00342     /** Speed value of the tuple (optional) */
00343     speed_t speed;
00344     /** Speed level value of the tuple (required) */
00345     speed_level_t speed_level;
00346 } hw_behavior_t;
00347
00348 /**
00349     * Defines policy for determining speed.
00350     */
00351 typedef long speed_policy_t;
00352
00353 /**
00354     * Defines policy for scheduling.
00355     */
00356 typedef long scheduling_policy_t;
00357
00358 /** @} */
00359
00360

```

```
00361 //=====
00362 // Error codes
00363 //-----
00364
00365 /** \addtogroup errors Error Codes
00366     * @{
00367     */
00368
00369 /** Error codes returned by API functions */
00370 typedef int  retval_t;
00371
00372 /** Feature unsupported by hardware */
00373 #define PWR_ARCH_UNSUPPORTED (-3)
00374
00375 /** Feature not implemented */
00376 #define PWR_UNIMPLEMENTED (-2)
00377
00378 /** Power API has not been initialized */
00379 #define PWR_UNINITIALIZED (-1)
00380
00381 /** Command executed successfully */
00382 #define PWR_OK (0)
00383
00384 /** Unspecified error */
00385 #define PWR_ERR (1)
00386
00387 /** Feature temporarily unavailable */
00388 #define PWR_UNAVAILABLE (2)
00389
00390 /** Request was denied */
00391 #define PWR_REQUEST_DENIED (4)
00392
00393 /** Unspecified error during initialization */
00394 #define PWR_INIT_ERR (5)
00395
00396 /** Unspecified error during finalization */
00397 #define PWR_FINAL_ERR (6)
00398
00399 /** Attempt to initialize API after it has been initialized */
00400 #define PWR_ALREADY_INITIALIZED (7)
00401
00402 /** Unspecified input / output error */
00403 #define PWR_IO_ERR (8)
00404
00405 /** Speed level not supported by hardware or API */
00406 #define PWR_UNSUPPORTED_SPEED_LEVEL (9)
00407
00408 /** Voltage not supported by hardware or API */
00409 #define PWR_UNSUPPORTED_VOLTAGE (10)
00410
00411 /** Feature is already set to minimum or maximum value */
00412 #define PWR_ALREADY_MINMAX (11)
00413
00414 /** Request denied, over energy budget */
00415 #define PWR_OVER_E_BUDGET (12) // reserved for future
00416
00417 /** Request denied, over power budget */
00418 #define PWR_OVER_P_BUDGET (13) // reserved for future
00419
00420 /** Request denied, over thermal budget */
00421 #define PWR_OVER_T_BUDGET (14) // reserved for future
00422
00423 /** Specified island does not exist */
00424 #define PWR_INVALID_ISLAND (15)
00425
00426 /** Unspecified error when changing voltage and/or frequency */
00427 #define PWR_DVFS_ERR (16)
00428
```

```

00429 /** Requested value overflow */
00430 #define PWR_OVERFLOW_ERR (17)
00431 /** @} */
00432
00433 //=====
00434 // High-Level Interface
00435 //-----
00436
00437 /** \addtogroup highlevel High-Level Interface
00438  * @{
00439  */
00440
00441 /**
00442  * Checks if the Power API has been initialized
00443  *
00444  * @param is_initialized <code>TRUE</code> if initialized, <code>FALSE</code>
00445  * otherwise
00446  * @retval #PWR_OK All cases
00447  */
00448 retval_t pwr_is_initialized(int* is_initialized);
00449
00450
00451 /**
00452  * Allocates resources used by the Power API
00453  *
00454  * Must be called before any other function in the Power API can be used.
00455  *
00456  * @param hw_behavior Relationship between voltage and speed on all voltage
00457  * islands
00458  * @param speed_policy Blanket policy for controlling voltage island speed
00459  * levels
00460  * @param scheduling_policy Blanket policy for runtime task scheduling
00461  *
00462  * @retval #PWR_OK Initialized with no errors
00463  * @retval #PWR_INIT_ERR Initialization failed
00464  * @retval #PWR_ALREADY_INITIALIZED API is initialized
00465  */
00466 retval_t pwr_initialize(const hw_behavior_t* hw_behavior,
00467                        const speed_policy_t* speed_policy,
00468                        const scheduling_policy_t* scheduling_policy);
00469
00470
00471 /**
00472  * Frees resources used by the Power API
00473  *
00474  * Must be called after all Power API functions have been called.
00475  *
00476  * @retval #PWR_OK All resources freed successfully
00477  * @retval #PWR_UNINITIALIZED Power API has not been initialized thus cannot
00478  * be finalized
00479  */
00480 retval_t pwr_finalize();
00481
00482 /**
00483  * This method is a quick fix to a strange problem caused by re-initializing
00484  * and finalizing the ecount environment multiple times. In particular,
00485  * if the ec_finalize method is called in-between measurments then for some
00486  * unknown reason measurments after ec_finalize always return zero.
00487  */
00488 retval_t pwr_ecount_finalize();
00489
00490 /**
00491  * Retrieve the currently active hardware behavior
00492  *
00493  * @param hw_behavior Pointer to active hardware behavior struct
00494  *
00495  * @retval #PWR_OK Active hardware behavior returned successfully

```

```

00496 * @retval #PWR_UNINITIALIZED Power API has not been initialized thus
00497 hardware
00498 * behavior is unavailable
00498 */
00499 retval_t pwr_hw_behavior(hw_behavior_t* hw_behavior);
00500
00501 /**
00502 * Change currently active hardware behavior
00503 *
00504 * @param hw_behavior Pointer to new active hardware behavior struct
00505 *
00506 * @retval #PWR_OK Active hardware behavior successfully changed
00507 * @retval #PWR_UNINITIALIZED Power API has not been initialized thus
00508 hardware
00509 * behavior cannot be changed
00509 */
00510 retval_t pwr_change_hw_behavior(hw_behavior_t* hw_behavior);
00511
00512
00513 /** @} */
00514
00515 //=====
00516 // Low-Level Interface
00517 //-----
00518
00519 /** \addtogroup lowlevel Low-Level Interface
00520 * @{
00521 */
00522 /**
00523 * The number of voltage islands controlled by the Power API
00524 *
00525 * @param num_islands Total voltage islands controlled by the Power API
00526 *
00527 * @retval #PWR_OK Number of voltage islands has been returned
00528 * @retval #PWR_UNINITIALIZED Power API has not been initialized
00529 */
00530 retval_t pwr_num_islands(long* num_islands);
00531
00532
00533 /**
00534 * Retrieve unique ID for all voltage islands
00535 *
00536 * Requires that <code>islands</code> points to at least
00537 * <code>num_islands*sizeof(island_id_t)</code> bytes of memory.
00538 *
00539 * @param islands An array of <code>island_id_t</code>
00540 *
00541 * @retval #PWR_OK Voltage island IDs have been returned
00542 * @retval #PWR_UNINITIALIZED Power API has not been initialized
00543 */
00544 retval_t pwr_islands(island_id_t* islands);
00545
00546
00547 /**
00548 * Number of discrete speed levels, not including power and clock gated
00549 states,
00550 * supported by a voltage island
00550 *
00551 * The slowest speed level is '<code>1</code>' and speed levels increase
00552 monotonically until the fastest speed level at
00553 '<code>num_speed_levels</code>'.
00554 *
00555 * @todo Map speed level to integer between 0 and 100 inclusive?
00556 *
00557 * @param island The ID of the island to get speed level count for
00558 * @param num_speed_levels The number of speed levels supported by

```

```

00559     <code>island</code>
00560     * @retval #PWR_OK Number of speed levels has been returned
00561     * @retval #PWR_INVALID_ISLAND The given ID does not correspond to a voltage
00562     * @retval #PWR_UNINITIALIZED Power API has not been initialized
00563     */
00564     retval_t pwr_num_speed_levels(const island_id_t island, long* num_speed_levels)
00565     ;
00566
00567 /**
00568  * The current speed level of a voltage island
00569  *
00570  * @param island The island to check speed level on
00571  * @param current_level The current speed level of <code>island</code>
00572  *
00573  * @retval #PWR_OK Speed level has been returned
00574  * @retval #PWR_INVALID_ISLAND The given ID does not correspond to a voltage
00575  * @retval #PWR_UNINITIALIZED Power API has not been initialized
00576  */
00577     retval_t pwr_current_speed_level(const island_id_t island, speed_level_t*
00578     current_level);
00579
00580 /**
00581  * Requests speed level on the given voltage island
00582  *
00583  * Requires <code>-1 <= new_level <= num_speed_levels</code>.
00584  * '<code>new_level == -1</code>' requests power gating.
00585  * '<code>new_level == 0</code>' requests clock gating.
00586  *
00587  * @param island The island to request speed level on
00588  * @param new_level The requested speed level
00589  *
00590  * @retval #PWR_OK Speed level has been changed
00591  * @retval #PWR_UNSUPPORTED_SPEED_LEVEL The requested speed level is not
00592  * supported
00593  * @retval #PWR_ALREADY_MINMAX Voltage island already at requested minimum or
00594  * maximum speed level
00595  * @retval #PWR_DVFS_ERR DVFS failed
00596  * @retval #PWR_INVALID_ISLAND The given ID does not correspond to a voltage
00597  * island.
00598  * @retval #PWR_OVER_E_BUDGET Request denied, over energy budget (reserved for
00599  * future use)
00600  * @retval #PWR_OVER_P_BUDGET Request denied, over power budget (reserved for
00601  * future use)
00602  * @retval #PWR_OVER_T_BUDGET Request denied, over thermal budget (reserved
00603  * for future use)
00604  * @retval #PWR_UNINITIALIZED Power API has not been initialized
00605  */
00606     retval_t pwr_request_speed_level(const island_id_t island, const speed_level_t
00607     new_level);
00608
00609 /**
00610  * Requests a speed level modification on the given island
00611  *
00612  * <code>delta</code> can be positive or negative.
00613  *
00614  * @param island The island to speed up or slow down
00615  * @param delta The number of speed levels to increment by
00616  * @param bottom The minimum legal speed level, can be one of
00617  * #PWR_POWER_GATE,
00618  * #PWR_CLOCK_GATE or #PWR_MIN_SPEED_LEVEL
00619  *
00620  * @todo Clarify behavior when an illegal speed level is requested

```

```

00614 * @todo Clarify semantics of speed levels when multiple hardware behaviors
00615 *       are used
00616 *
00617 * @return #PWR_OK Speed level was successfully modified
00618 * @retval #PWR_UNSUPPORTED_SPEED_LEVEL The requested speed level is not
supported
00619 * @retval #PWR_ALREADY_MINMAX Voltage island already at requested minimum or
maximum speed level
00620 * @retval #PWR_DVFS_ERR DVFS failed
00621 * @retval #PWR_INVALID_ISLAND The given ID does not correspond to a voltage
island.
00622 * @retval #PWR_OVER_E_BUDGET Request denied, over energy budget (reserved for
future use)
00623 * @retval #PWR_OVER_P_BUDGET Request denied, over power budget (reserved for
future use)
00624 * @retval #PWR_OVER_T_BUDGET Request denied, over thermal budget (reserved
for future use)
00625 * @retval #PWR_UNINITIALIZED Power API has not been initialized
00626 */
00627 retval_t pwr_modify_speed_level(const island_id_t island,
00628                                const int delta,
00629                                const speed_level_t bottom);
00630
00631 /**
00632 * Calculates the cost of switching speed levels
00633 *
00634 * @todo Decide on unit for agility. Currently ns because of cpufreq default.
00635 *       If agility in cycles, at what speed level? Full speed?
00636 *       <code>to_level</code>?
00637 *
00638 * @param island The island of interest
00639 * @param from_level Starting speed level
00640 * @param to_level Finishing speed level
00641 * @param best_case Minimum cost in ns
00642 * @param worst_case Maximum cost in ns
00643 *
00644 * @retval #PWR_OK Agility successfully returned
00645 * @retval #PWR_ERR Agility not returned
00646 */
00647 retval_t pwr_agility(const island_id_t island,
00648                     const speed_t from_level,
00649                     const speed_t to_level,
00650                     agility_t* best_case,
00651                     agility_t* worst_case);
00652
00653 /**
00654 * Requests a voltage level modification of the given island
00655 *
00656 * <code>delta</code> can be positive or negative.
00657 *
00658 * @todo Should we have a corresponding function to directly set voltage as
00659 *       we do with speed level?
00660 *
00661 * @param island The island to change voltage on
00662 * @param delta The number of voltage levels to modify by
00663 *
00664 * @retval #PWR_OK Voltage level successfully modified
00665 * @retval #PWR_ARCH_UNSUPPORTED Voltage level cannot be modified
00666 * @retval #PWR_UNSUPPORTED_SPEED_LEVEL The requested speed level is not
supported
00667 * @retval #PWR_ALREADY_MINMAX Voltage island already at requested minimum or
maximum voltage
00668 * @retval #PWR_DVFS_ERR DVFS failed
00669 * @retval #PWR_INVALID_ISLAND The given ID does not correspond to a voltage
island.
00670 * @retval #PWR_OVER_E_BUDGET Request denied, over energy budget (reserved for
future use)
00671 * @retval #PWR_OVER_P_BUDGET Request denied, over power budget (reserved for

```

```

future use)
00672  * @retval #PWR_OVER_T_BUDGET Request denied, over thermal budget (reserved
for future use)
00673  * @retval #PWR_UNINITIALIZED Power API has not been initialized
00674  * @retval #PWR_ERR Voltage level not successfully modified
00675  */
00676 retval_t pwr_modify_voltage(const island_id_t island, const int delta);
00677
00678 /**
00679  * Value of monotonically increasing energy counter on the given island.
00680  * Total energy consumed (J) since some undefined starting point is
00681  * <code>e_j + e_uj*10E6</code>. Elapsed time (sec) since some undefined
00682  * starting point is <code>t_sec + t_nsec*10E9</code>.
00683  *
00684  * @param island The island to measure energy consumption on
00685  * @param e_j Energy consumed since some unspecified starting point on
00686  * <code>island</code>, joule component
00687  * @param e_uj Energy consumed since some undefined starting point on
00688  * <code>island</code>, microjoule component
00689  * @param t_sec Time since some unspecified starting point, seconds component
00690  * @param t_nsec Time since some unspecified starting point, nanoseconds
component
00691  *
00692  * @retval #PWR_OK Energy counter successfully returned
00693  * @retval #PWR_ERR Energy counter not returned
00694  * @retval #PWR_OVERFLOW_ERR Energy counter overflowed, inaccurate results
provided
00695  */
00696 retval_t pwr_energy_counter(const island_id_t island,
00697                             energy_t* e_j,
00698                             energy_t* e_uj,
00699                             timestamp_t* t_sec,
00700                             timestamp_t* t_nsec);
00701 /** @} */
00702 #ifdef __cplusplus
00703 }
00704 #endif
00705 #endif

```